

Biostat 212a Homework 4

Due Mar. 4, 2025 @ 11:59PM

Jiaye Tian UID: 306541095

2025-03-04

Table of contents

0.1	ISL Exercise 8.4.3 (10pts)	1
0.2	ISL Exercise 8.4.4 (10pts)	2
0.3	ISL Exercise 8.4.5 (10pts)	4
0.4	ISL Lab 8.3. Boston data set (30pts)	5
0.5	ISL Lab 8.3 Carseats data set (30pts)	14


0.1 ISL Exercise 8.4.3 (10pts)

```
p = seq(0, 1, 0.01)
gini = p * (1 - p) * 2
entropy = -(p * log(p) + (1 - p) * log(1 - p))
class.err = 1 - pmax(p, 1 - p)

matplot(p, cbind(gini, entropy, class.err), type = "l",
        col = c("red", "green", "blue"),
        main = "Comparison of Gini Index, Entropy, and Classification Error")


legend("topright", legend=c("Gini", "Entropy", "Class Error"), pch=19,
      col=c("red", "green", "blue"))
```

Comparison of Gini Index, Entropy, and Classification Error



0.2 ISL Exercise 8.4.4 (10pts)

(a)



(b)

x_2	2	2.49	
	1	-1.06	0.21
		-1.8	0.63
		0	1

0.3 ISL Exercise 8.4.5 (10pts)

Majority Vote \rightarrow Red Averaging Probability \rightarrow Green

If $P(\text{Red} | X) > 0.5 \Rightarrow$ Classifier predicts Red.

If $P(\text{Red} | X) < 0.5 \Rightarrow$ Classifier predicts Green

$P(\text{Red} | X) \leq 0.5$: 0.1, 0.15, 0.2, 0.2 \Rightarrow 4 classifiers predict Green.

$P(\text{Red} | X) > 0.5$: 0.55, 0.6, 0.6, 0.65, 0.7, 0.75 \Rightarrow 6 classifiers predict Red.

Thus, the majority (6 out of 10) Vote \Rightarrow Red. \square

$$\begin{aligned} P_{\text{avg}}(\text{Red} | X) &= \frac{0.1 + 0.15 + 0.2 + 0.2 + 0.55 + 0.6 + 0.6 + 0.65 + 0.7 + 0.75}{10} \\ &= 0.45 \end{aligned}$$

$$P_{\text{avg}}(\text{Green} | X) = 1 - 0.45 = 0.55 \Rightarrow \text{Green}.$$

0.4 ISL Lab 8.3. Boston data set (30pts)

Follow the machine learning workflow to train regression tree, random forest, and boosting methods for predicting `medv`. Evaluate out-of-sample performance on a test set.

```
# Load necessary libraries
library(MASS)           # Boston dataset
library(tree)           # Regression Tree
library(randomForest)   # Random Forest
library(gbm)            # Boosting (Gradient Boosting)
# library(ISLR2)
# attach(Carseats)

# Set seed for reproducibility
set.seed(1)

# Split the dataset into training (50%) and testing (50%)
train <- sample(1:nrow(Boston), nrow(Boston) / 2)
Boston.test <- Boston[-train, "medv"] # Extract true test set values
```

1. Regression tree

```
# Train Regression Tree
tree.Boston <- tree(medv ~ ., Boston, subset = train)
summary(tree.Boston)
```

Regression tree:

```
tree(formula = medv ~ ., data = Boston, subset = train)
```

Variables actually used in tree construction:

```
[1] "rm"      "lstat" "crim"  "age"
```


Number of terminal nodes: 7

Residual mean deviance: 10.38 = 2555 / 246

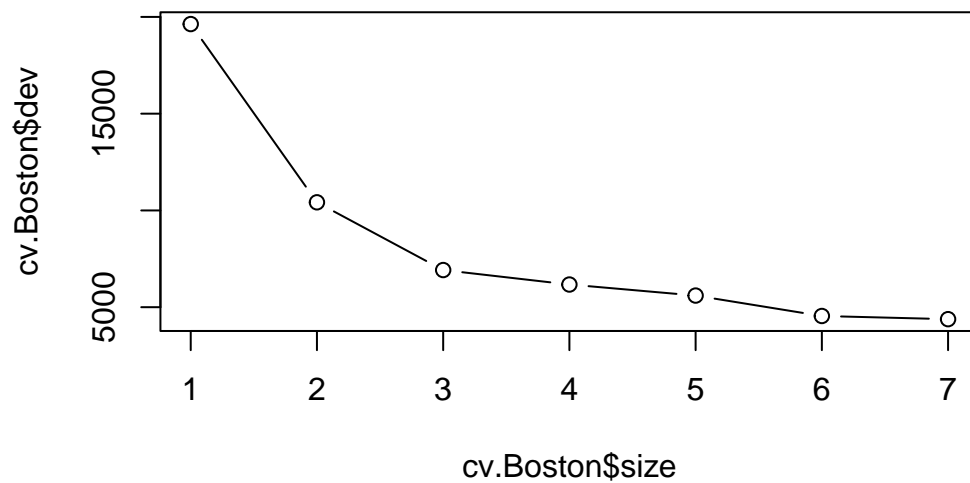
Distribution of residuals:

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-10.1800	-1.7770	-0.1775	0.0000	1.9230	16.5800


```
# Plot the tree
plot(tree.Boston)
text(tree.Boston)
```




```
# Cross-validation for pruning
cv.Boston <- cv.tree(tree.Boston)
plot(cv.Boston$size, cv.Boston$dev, type = "b")
```



```
# Prune the tree with best = 5 (per original code)
prune.Boston <- prune.tree(tree.Boston, best = 5)
plot(prune.Boston)
text(prune.Boston, pretty = 5)
```



```
# Predict on test set and compute MSE
yhat.tree <- predict(tree.Boston, newdata = Boston[-train, ])
Boston.test <- Boston[-train, "medv"]
plot(yhat.tree, Boston.test)
abline(0, 1)
```

```
# Compute the test set MSE
tree.mse <- mean((yhat.tree - Boston.test)^2)
tree.mse
```

```
[1] 35.28688
```

In other words, the test set MSE associated with the regression tree is 35.29. The square root of the MSE is therefore around 5.941, indicating that this model leads to test predictions that are (on average) within approximately \$5,941 of the true median home value for the census tract.

2. RandomForest

```
# Train Bagging model (mtry = 12, per original code)
set.seed(1)
bag.Boston <- randomForest(medv ~ ., data = Boston,
                           subset = train, mtry = 12,
                           importance = TRUE)

# Test set prediction & MSE
yhat.bag <- predict(bag.Boston, newdata = Boston[-train, ])
bag.mse <- mean((yhat.bag - Boston.test)^2)
bag.mse
```

```
[1] 23.38773
```

```
# Train Bagging with ntree = 25 (per original code)
bag.Boston <- randomForest(medv ~ ., data = Boston,
                           subset = train, mtry = 12, ntree = 25)
yhat.bag <- predict(bag.Boston, newdata = Boston[-train, ])
bag.mse2 <- mean((yhat.bag - Boston.test)^2)
bag.mse2
```

```
[1] 25.19144
```

```
# Train Random Forest (mtry = 6, per original code)
set.seed(1)
rf.Boston <- randomForest(medv ~ ., data = Boston,
                          subset = train, mtry = 6, importance = TRUE)

# Test set prediction & MSE
yhat.rf <- predict(rf.Boston, newdata = Boston[-train, ])
rf.mse <- mean((yhat.rf - Boston.test)^2)
rf.mse
```


```
[1] 19.62021
```

```
# Feature importance & visualization
importance(rf.Boston)
```

	%IncMSE	IncNodePurity
crim	16.697017	1076.08786
zn	3.625784	88.35342
indus	4.968621	609.53356
chas	1.061432	52.21793
nox	13.518179	709.87339
rm	32.343305	7857.65451
age	13.272498	612.21424
dis	9.032477	714.94674
rad	2.878434	95.80598
tax	9.118801	364.92479
ptratio	8.467062	823.93341
black	7.579482	275.62272
lstat	27.129817	6027.63740

```
varImpPlot(rf.Boston)
```


rf.Boston




3. Boosting

```
# Train Boosting model (default shrinkage = 0.1)
set.seed(1)
boost.Boston <- gbm(medv ~ ., data = Boston[train, ],
distribution = "gaussian", n.trees = 5000, interaction.depth = 4)

# Visualize variable importance
plot(boost.Boston, i = "rm")
```



```
plot(boost.Boston, i = "lstat")
```



```
# Test set prediction & MSE
yhat.boost <- predict(boost.Boston, newdata = Boston[-train, ], n.trees = 5000)
boost.mse <- mean((yhat.boost - Boston.test)^2)
boost.mse
```

```
[1] 18.84709
```

```
# Train Boosting with shrinkage = 0.2 (per original code)
boost.Boston <- gbm(medv ~ ., data = Boston[train, ],
  distribution = "gaussian", n.trees = 5000,
  interaction.depth = 4, shrinkage = 0.2, verbose = F)

# Test set prediction & MSE
yhat.boost <- predict(boost.Boston, newdata = Boston[-train, ], n.trees = 5000)
boost.mse2 <- mean((yhat.boost - Boston.test)^2)
boost.mse2
```


```
[1] 18.33455
```

```
# Compare MSE results
mse_results <- data.frame(
  Model = c("Regression Tree", "Bagging (mtry=12, ntree=500)",
    "Bagging (mtry=12, ntree=25)", "Random Forest",
    "Boosting (shrinkage=0.1)", "Boosting (shrinkage=0.2)"),
  MSE = c(tree.mse, bag.mse, bag.mse2, rf.mse, boost.mse, boost.mse2)
)
print(mse_results)
```

	Model	MSE
1	Regression Tree	35.28688
2	Bagging (mtry=12, ntree=500)	23.38773
3	Bagging (mtry=12, ntree=25)	25.19144
4	Random Forest	19.62021
5	Boosting (shrinkage=0.1)	18.84709
6	Boosting (shrinkage=0.2)	18.33455

```
# Plot MSE comparison
barplot(mse_results$MSE, names.arg = mse_results$Model,
  col = c("red", "blue", "blue", "green", "orange", "orange"),
  main = "MSE Comparison", ylab = "Mean Squared Error", las = 2, cex.names = 0.8)
```

MSE Comparison



0.5 ISL Lab 8.3 Carseats data set (30pts)

Follow the machine learning workflow to train classification tree, random forest, and boosting methods for classifying Sales ≤ 8 versus Sales > 8 . Evaluate out-of-sample performance on a test set.

```
# Load required libraries
library(tree)           # Classification Tree
library(randomForest)   # Random Forest
library(gbm)            # Gradient Boosting
library(caret)          # Confusion Matrix
library(ISLR2)

# Set seed for reproducibility
set.seed(2)

# Convert Sales into a binary variable
Carseats$High <- factor(ifelse(Carseats$Sales > 8, "Yes", "No"))

# Remove the Sales column
Carseats <- subset(Carseats, select = -Sales)
```

```
# Train-test split (50-50)
train <- sample(1:nrow(Carseats), nrow(Carseats) / 2)
Carseats.train <- Carseats[train, ]
Carseats.test <- Carseats[-train, ]
High.test <- Carseats.test$High

# ---- Classification Tree ----
tree.carseats <- tree(High ~ ., data = Carseats.train)
summary(tree.carseats)
```

Classification tree:

```
tree(formula = High ~ ., data = Carseats.train)
```

Variables actually used in tree construction:

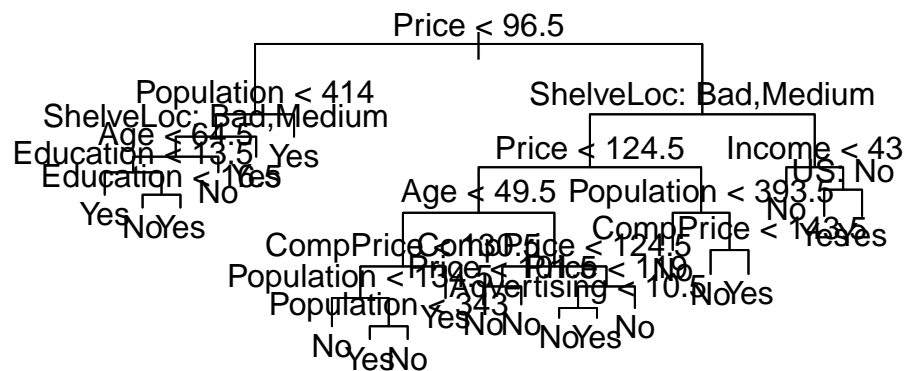
```
[1] "Price"          "Population"    "ShelveLoc"     "Age"           "Education"
[6] "CompPrice"      "Advertising"   "Income"        "US"
```

Number of terminal nodes: 21

Residual mean deviance: 0.5543 = 99.22 / 179

Misclassification error rate: 0.115 = 23 / 200

```
plot(tree.carseats)
text(tree.carseats, pretty = 0)
```



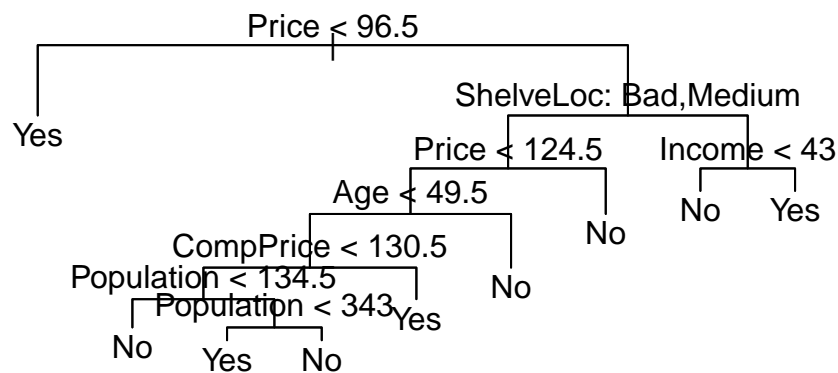
```

# Test set prediction
tree.pred <- predict(tree.carseats, Carseats.test, type = "class")
confusion_matrix_tree <- table(tree.pred, High.test)
tree_accuracy <- sum(diag(confusion_matrix_tree)) / sum(confusion_matrix_tree)

# ---- Pruning ----
set.seed(7)
cv.carseats <- cv.tree(tree.carseats, FUN = prune.misclass)
prune_size <- cv.carseats$size[which.min(cv.carseats$dev)] # Select pruning size with minimum
prune.carseats <- prune.misclass(tree.carseats, best = prune_size)

plot(prune.carseats)
text(prune.carseats, pretty = 0)

```



```

# Evaluate pruned tree
tree.pred.pruned <- predict(prune.carseats, Carseats.test, type = "class")
confusion_matrix_pruned <- table(tree.pred.pruned, High.test)
pruned_accuracy <- sum(diag(confusion_matrix_pruned)) / sum(confusion_matrix_pruned)


# ---- Random Forest ----
p <- ncol(Carseats.train) - 1 # Number of features excluding 'High'

```



```
# Tune mtry using tuneRF()
set.seed(1)
best_mtry <- tuneRF(Carseats.train[-ncol(Carseats.train)], Carseats.train$High,
                    stepFactor = 1.5, improve = 0.01, trace = FALSE)
```

```
-0.245283 0.01
-0.09433962 0.01
```



```
mtry_best <- best_mtry[which.min(best_mtry[,2]), 1]

# Train Random Forest
set.seed(1)
rf.carseats <- randomForest(High ~ ., data = Carseats.train, ntree = 500, mtry = mtry_best,

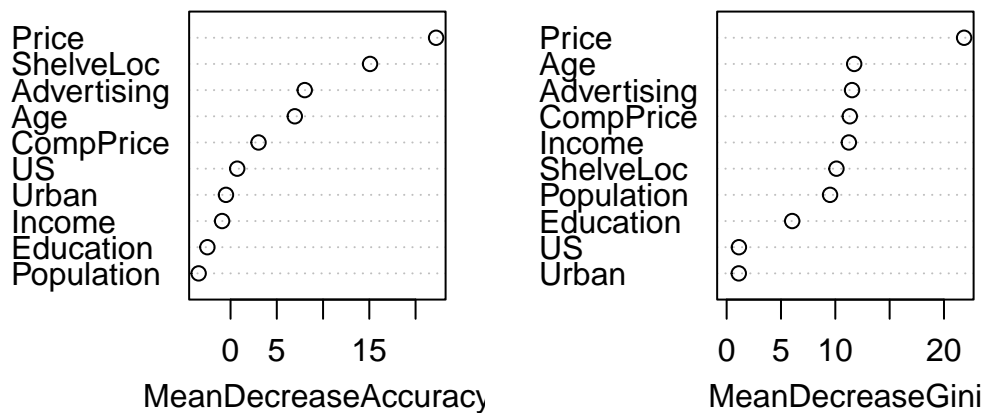
# Test set prediction
rf.pred <- predict(rf.carseats, Carseats.test)
confusion_matrix_rf <- table(rf.pred, High.test)
rf_accuracy <- sum(diag(confusion_matrix_rf)) / sum(confusion_matrix_rf)

# Variable Importance
importance(rf.carseats)
```

	No	Yes	MeanDecreaseAccuracy	MeanDecreaseGini
CompPrice	2.1304913	2.21583500	3.0233624	11.332341
Income	-1.2187382	0.08332858	-0.9127061	11.247612
Advertising	2.7107740	9.28221142	8.0200032	11.539777
Population	-2.5279538	-1.96748887	-3.4496266	9.515984
Price	17.6102844	16.18029714	22.2205678	21.853751
ShelveLoc	11.8799961	11.50873822	15.0797611	10.106813
Age	3.7038805	6.83255965	6.9440472	11.738663
Education	-2.9218854	-0.16331173	-2.5039109	6.031744
Urban	-0.9337001	0.16405241	-0.4965739	1.114652
US	-0.4330561	1.48555372	0.7195542	1.127466

```
varImpPlot(rf.carseats)
```

rf.carseats



```
# ---- Boosting ----
boost.train <- Carseats.train
boost.test <- Carseats.test
boost.train$High <- ifelse(boost.train$High == "Yes", 1, 0) # Convert to 0/1 for Boosting

set.seed(1)
boost.carseats <- gbm(High ~ ., data = boost.train, distribution = "bernoulli",
                      n.trees = 5000, interaction.depth = 4, shrinkage = 0.01, verbose = FALSE)
```


```

# Prediction
boost.prob <- predict(boost.carseats, newdata = boost.test, n.trees = 5000, type = "response")
boost.pred <- factor(ifelse(boost.prob > 0.5, "Yes", "No"), levels = levels(High.test))

confusion_matrix_boost <- table(boost.pred, High.test)
boost_accuracy <- sum(diag(confusion_matrix_boost)) / sum(confusion_matrix_boost)

# Variable Importance
summary(boost.carseats)

```



	var	rel.inf
Price	Price	24.3344414
CompPrice	CompPrice	14.3555860
ShelveLoc	ShelveLoc	13.4625168
Age	Age	12.5211889
Advertising	Advertising	11.2205000
Income	Income	11.1522623
Population	Population	8.0522931
Education	Education	3.4101404
Urban	Urban	0.8634069
US	US	0.6276643

```
# Test Boosting with different shrinkage (0.2)
set.seed(1)
boost.carseats2 <- gbm(High ~ ., data = boost.train, distribution = "bernoulli",
                      n.trees = 5000, interaction.depth = 4, shrinkage = 0.2, verbose = FALSE)

boost.prob2 <- predict(boost.carseats2, newdata = boost.test, n.trees = 5000, type = "response")
boost.pred2 <- factor(ifelse(boost.prob2 > 0.5, "Yes", "No"), levels = levels(High.test))

confusion_matrix_boost2 <- table(boost.pred2, High.test)
boost_accuracy2 <- sum(diag(confusion_matrix_boost2)) / sum(confusion_matrix_boost2)

# ---- Confusion Matrix Visualization ----
cat("Classification Tree:\n")
```

Classification Tree:

```
print(confusionMatrix(as.factor(tree.pred), High.test))
```

Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	104	33
Yes	13	50

Accuracy : 0.77
 95% CI : (0.7054, 0.8264)
 No Information Rate : 0.585
 P-Value [Acc > NIR] : 2.938e-08

Kappa : 0.5091

McNemar's Test P-Value : 0.005088

Sensitivity : 0.8889
 Specificity : 0.6024
 Pos Pred Value : 0.7591
 Neg Pred Value : 0.7937
 Prevalence : 0.5850
 Detection Rate : 0.5200
 Detection Prevalence : 0.6850

Balanced Accuracy : 0.7456

'Positive' Class : No

```
cat("\nPruned Tree:\n")
```

Pruned Tree:

```
print(confusionMatrix(as.factor(tree.pred.pruned), High.test))
```

Confusion Matrix and Statistics

Reference

Prediction No Yes

No 97 25

Yes 20 58

Accuracy : 0.775

95% CI : (0.7108, 0.8309)

No Information Rate : 0.585

P-Value [Acc > NIR] : 1.206e-08

Kappa : 0.5325

McNemar's Test P-Value : 0.551

Sensitivity : 0.8291

Specificity : 0.6988

Pos Pred Value : 0.7951

Neg Pred Value : 0.7436

Prevalence : 0.5850

Detection Rate : 0.4850

Detection Prevalence : 0.6100

Balanced Accuracy : 0.7639

'Positive' Class : No

```
cat("\nRandom Forest:\n")
```

Random Forest:

```
print(confusionMatrix(as.factor(rf.pred), High.test))
```

Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	110	24
Yes	7	59

Accuracy : 0.845
95% CI : (0.7873, 0.8922)
No Information Rate : 0.585
P-Value [Acc > NIR] : 1.939e-15

Kappa : 0.671

McNemar's Test P-Value : 0.004057

Sensitivity : 0.9402
Specificity : 0.7108
Pos Pred Value : 0.8209
Neg Pred Value : 0.8939
Prevalence : 0.5850
Detection Rate : 0.5500
Detection Prevalence : 0.6700
Balanced Accuracy : 0.8255

'Positive' Class : No

```
cat("\nBoosting (shrinkage=0.01):\n")
```

Boosting (shrinkage=0.01):

```
print(confusionMatrix(as.factor(boost.pred), High.test))
```

Confusion Matrix and Statistics

```

      Reference
Prediction No Yes
      No  109  16
      Yes   8  67

      Accuracy : 0.88
      95% CI : (0.8267, 0.9216)
      No Information Rate : 0.585
      P-Value [Acc > NIR] : <2e-16

      Kappa : 0.7493

      Mcnemar's Test P-Value : 0.153

      Sensitivity : 0.9316
      Specificity : 0.8072
      Pos Pred Value : 0.8720
      Neg Pred Value : 0.8933
      Prevalence : 0.5850
      Detection Rate : 0.5450
      Detection Prevalence : 0.6250
      Balanced Accuracy : 0.8694

      'Positive' Class : No
```

```
cat("\nBoosting (shrinkage=0.2):\n")
```

Boosting (shrinkage=0.2):

```
print(confusionMatrix(as.factor(boost.pred2), High.test))
```

Confusion Matrix and Statistics

```

      Reference
```

Prediction	No	Yes
No	111	36
Yes	6	47

Accuracy : 0.79
 95% CI : (0.7269, 0.8443)
 No Information Rate : 0.585
 P-Value [Acc > NIR] : 7.046e-10

Kappa : 0.5435

McNemar's Test P-Value : 7.648e-06

Sensitivity : 0.9487
 Specificity : 0.5663
 Pos Pred Value : 0.7551
 Neg Pred Value : 0.8868
 Prevalence : 0.5850
 Detection Rate : 0.5550
 Detection Prevalence : 0.7350
 Balanced Accuracy : 0.7575

'Positive' Class : No

```

# ---- Accuracy Comparison ----
accuracy_results <- data.frame(
  Model = c("Classification Tree", "Pruned Tree", "Random Forest", "Boosting (0.01)", "Boosting (0.2)"),
  Accuracy = c(tree_accuracy, pruned_accuracy, rf_accuracy, boost_accuracy, boost_accuracy2)
)

print(accuracy_results)

```

	Model	Accuracy
1	Classification Tree	0.770
2	Pruned Tree	0.775
3	Random Forest	0.845
4	Boosting (0.01)	0.880
5	Boosting (0.2)	0.790


```
barplot(accuracy_results$Accuracy, names.arg = accuracy_results$Model,  
        col = c("red", "blue", "green", "orange", "purple"),  
        main = "Accuracy Comparison", ylab = "Accuracy", las = 2, cex.names = 0.8)
```

